

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torchsummary import summary
import matplotlib.pyplot as plt
```

```
from sklearn.metrics import confusion_matrix
import seaborn as sn
import pandas as pd
import numpy as np
```

```
# Set the path where you want to save the checkpoint
path = "/content/drive/MyDrive/my_checkpoint_cosineLR_v1.pth"
```

```
# Define the device to use for training
device = torch.device("cuda:0" if torch.cuda.is_available() else
"cpu")
```

```
# Define the CNN architecture
```

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=32,
kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(32)
        self.relu1 = nn.ReLU()
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=64,
kernel_size=3, padding=1)
        self.bn2 = nn.BatchNorm2d(64)
        self.relu2 = nn.ReLU()
        self.maxpool1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv3 = nn.Conv2d(in_channels=64, out_channels=128,
kernel_size=3, padding=1)
        self.bn3 = nn.BatchNorm2d(128)
        self.relu3 = nn.ReLU()
        self.conv4 = nn.Conv2d(in_channels=128, out_channels=256,
kernel_size=3, padding=1)
        self.bn4 = nn.BatchNorm2d(256)
        self.relu4 = nn.ReLU()
        self.maxpool2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv5 = nn.Conv2d(in_channels=256, out_channels=512,
kernel_size=3, padding=1)
```

```

        self.bn5 = nn.BatchNorm2d(512)
        self.relu5 = nn.ReLU()
        self.conv6 = nn.Conv2d(in_channels=512, out_channels=1024,
kernel_size=3, padding=1)
        self.bn6 = nn.BatchNorm2d(1024)
        self.relu6 = nn.ReLU()
        self.maxpool3 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.Linear(in_features=4*4*1024, out_features=512)
        self.relu7 = nn.ReLU()
        self.dropout1 = nn.Dropout(p=0.5)
        self.fc2 = nn.Linear(in_features=512, out_features=10)

```

```

def forward(self, x):
    x = self.conv1(x)
    x = self.bn1(x)
    x = self.relu1(x)
    x = self.conv2(x)
    x = self.bn2(x)
    x = self.relu2(x)
    x = self.maxpool1(x)
    x = self.conv3(x)
    x = self.bn3(x)
    x = self.relu3(x)
    x = self.conv4(x)
    x = self.bn4(x)
    x = self.relu4(x)
    x = self.maxpool2(x)
    x = self.conv5(x)
    x = self.bn5(x)
    x = self.relu5(x)
    x = self.conv6(x)
    x = self.bn6(x)
    x = self.relu6(x)
    x = self.maxpool3(x)
    x = x.view(x.size(0), -1)
    x = self.fc1(x)
    x = self.relu7(x)
    x = self.dropout1(x)
    x = self.fc2(x)
    return x

```

```

net = Net()
net.cuda()

```

```

Net(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu1): ReLU()
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1),

```

```
padding=(1, 1))
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu2): ReLU()
    (maxpool1): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
    (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu3): ReLU()
    (conv4): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (bn4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu4): ReLU()
    (maxpool2): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
    (conv5): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (bn5): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu5): ReLU()
    (conv6): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (bn6): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu6): ReLU()
    (maxpool3): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
    (fc1): Linear(in_features=16384, out_features=512, bias=True)
    (relu7): ReLU()
    (dropout1): Dropout(p=0.5, inplace=False)
    (fc2): Linear(in_features=512, out_features=10, bias=True)
)

summary(net, (3, 32, 32))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 32, 32]	896
BatchNorm2d-2	[-1, 32, 32, 32]	64
ReLU-3	[-1, 32, 32, 32]	0
Conv2d-4	[-1, 64, 32, 32]	18,496
BatchNorm2d-5	[-1, 64, 32, 32]	128
ReLU-6	[-1, 64, 32, 32]	0
MaxPool2d-7	[-1, 64, 16, 16]	0
Conv2d-8	[-1, 128, 16, 16]	73,856
BatchNorm2d-9	[-1, 128, 16, 16]	256
ReLU-10	[-1, 128, 16, 16]	0

Conv2d-11	[-1, 256, 16, 16]	295,168
BatchNorm2d-12	[-1, 256, 16, 16]	512
ReLU-13	[-1, 256, 16, 16]	0
MaxPool2d-14	[-1, 256, 8, 8]	0
Conv2d-15	[-1, 512, 8, 8]	1,180,160
BatchNorm2d-16	[-1, 512, 8, 8]	1,024
ReLU-17	[-1, 512, 8, 8]	0
Conv2d-18	[-1, 1024, 8, 8]	4,719,616
BatchNorm2d-19	[-1, 1024, 8, 8]	2,048
ReLU-20	[-1, 1024, 8, 8]	0
MaxPool2d-21	[-1, 1024, 4, 4]	0
Linear-22	[-1, 512]	8,389,120
ReLU-23	[-1, 512]	0
Dropout-24	[-1, 512]	0
Linear-25	[-1, 10]	5,130

```
=====
Total params: 14,686,474
Trainable params: 14,686,474
Non-trainable params: 0
-----
```

```
Input size (MB): 0.01
Forward/backward pass size (MB): 7.14
Params size (MB): 56.02
Estimated Total Size (MB): 63.17
-----
```

```
# Define the transformations for the dataset
transform_train = transforms.Compose(
    [transforms.Resize((32,32)), #resizes the image so it can be
perfect for our model.
    transforms.RandomHorizontalFlip(), # FLips the image w.r.t
horizontal axis
    transforms.RandomRotation(10), #Rotates the image to a
specified angel
    transforms.RandomAffine(0, shear=10, scale=(0.8,1.2)), #Performs
actions like zooms, change shear angles.
    transforms.RandomCrop(32, padding=4),
    transforms.ColorJitter(brightness=0.2, contrast=0.2,
saturation=0.2),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

transform_test = transforms.Compose(
    [transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

# Download and load the CIFAR-10 dataset
train_dataset = torchvision.datasets.CIFAR10(root='./data',
train=True,
```

```

                                download=True,
transform=transform_train)
train_loader = torch.utils.data.DataLoader(train_dataset,
batch_size=32,
                                shuffle=True, num_workers=2)

test_dataset = torchvision.datasets.CIFAR10(root='./data',
train=False,
                                download=True,
transform=transform_test)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=32,
                                shuffle=False, num_workers=2)

```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to
./data/cifar-10-python.tar.gz

100%|██████████| 170498071/170498071 [00:02<00:00, 76882319.51it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified

```

# Define the loss function and the optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.005, momentum=0.9,
weight_decay=5e-4)
#scheduler = ReduceLR0nPlateau(optimizer, mode='min', factor=0.1,
patience=3, verbose=True)
scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=50,
eta_min=1e-8)

```

```

# Set the device to use (GPU or CPU)
device = torch.device("cuda:0" if torch.cuda.is_available() else
"cpu")

```

```

# initialize lists to store loss and accuracy
train_loss = []
train_acc = []
v_loss = []
v_acc = []
learn_rate = []

```

```

num_epochs = 50
# train the model
for epoch in range(num_epochs):
    running_loss = 0.0
    running_corrects = 0
    for images, labels in train_loader:
        # move data to device
        images, labels = images.to(device), labels.to(device)
        # zero the parameter gradients
        optimizer.zero_grad()

```

```

    # forward pass
    outputs = net(images)
    loss = criterion(outputs, labels)
    # backward pass and optimize
    loss.backward()
    optimizer.step()
    # calculate running loss and accuracy
    running_loss += loss.item() * images.size(0)
    _, preds = torch.max(outputs, 1)
    running_corrects += torch.sum(preds == labels.data)
    epoch_loss = running_loss / len(train_dataset)
    epoch_acc = 100 * running_corrects.double() / len(train_dataset)
    train_loss.append(epoch_loss)
    train_acc.append(epoch_acc)
    # print statistics
    print('Epoch [{}/{}], Training Loss: {:.4f}, Training Accuracy:
{:.4f} %'.format(epoch+1, num_epochs, epoch_loss, epoch_acc))

# Validate the model
net.eval()
val_loss = 0.0
val_correct = 0
total = 0

with torch.no_grad():
    for data in test_loader:
        inputs, labels = data[0].to(device), data[1].to(device)
        outputs = net(inputs)
        loss = criterion(outputs, labels)

        val_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        val_correct += (predicted == labels).sum().item()

    val_acc = 100 * val_correct / total
    val_loss /= len(test_loader)
    v_acc.append(val_acc)
    v_loss.append(val_loss)

# Print the training and validation metrics for each epoch
print('Test Epoch [{}/{}], Validation Loss: {:.4f}, Validation
Accuracy: {:.2f}%'.
    .format(epoch+1, num_epochs, val_loss, val_acc))

# Update the learning rate scheduler
# scheduler.step(val_loss)
scheduler.step()
epoch_lr = optimizer.param_groups[0]['lr']

```

```
print("Epoch:", epoch+1, "Learning rate:", epoch_lr )  
learn_rate.append(epoch_lr)
```

Epoch [1/50], Training Loss: 2.0764, Training Accuracy: 20.7200 %
Test Epoch [1/50], Validation Loss: 1.8293, Validation Accuracy: 33.61%
Epoch: 1 Learning rate: 0.0049950668309370365
Epoch [2/50], Training Loss: 1.5277, Training Accuracy: 43.4580 %
Test Epoch [2/50], Validation Loss: 1.2671, Validation Accuracy: 54.79%
Epoch: 2 Learning rate: 0.004980286792712688
Epoch [3/50], Training Loss: 1.1756, Training Accuracy: 57.6320 %
Test Epoch [3/50], Validation Loss: 0.8849, Validation Accuracy: 68.19%
Epoch: 3 Learning rate: 0.004955718215385468
Epoch [4/50], Training Loss: 0.9834, Training Accuracy: 65.4460 %
Test Epoch [4/50], Validation Loss: 0.7823, Validation Accuracy: 72.45%
Epoch: 4 Learning rate: 0.004921458059905771
Epoch [5/50], Training Loss: 0.8572, Training Accuracy: 70.0420 %
Test Epoch [5/50], Validation Loss: 0.6847, Validation Accuracy: 75.89%
Epoch: 5 Learning rate: 0.004877641535455302
Epoch [6/50], Training Loss: 0.7656, Training Accuracy: 73.2780 %
Test Epoch [6/50], Validation Loss: 0.5794, Validation Accuracy: 80.40%
Epoch: 6 Learning rate: 0.004824441565838198
Epoch [7/50], Training Loss: 0.6979, Training Accuracy: 75.7200 %
Test Epoch [7/50], Validation Loss: 0.6195, Validation Accuracy: 78.46%
Epoch: 7 Learning rate: 0.004762068107029786
Epoch [8/50], Training Loss: 0.6493, Training Accuracy: 77.4260 %
Test Epoch [8/50], Validation Loss: 0.5301, Validation Accuracy: 82.39%
Epoch: 8 Learning rate: 0.004690767318576258
Epoch [9/50], Training Loss: 0.6088, Training Accuracy: 78.9640 %
Test Epoch [9/50], Validation Loss: 0.5176, Validation Accuracy: 82.20%
Epoch: 9 Learning rate: 0.0046108205921154095
Epoch [10/50], Training Loss: 0.5670, Training Accuracy: 80.2120 %
Test Epoch [10/50], Validation Loss: 0.4583, Validation Accuracy: 84.41%
Epoch: 10 Learning rate: 0.004522543440852396
Epoch [11/50], Training Loss: 0.5326, Training Accuracy: 81.5060 %
Test Epoch [11/50], Validation Loss: 0.4503, Validation Accuracy: 84.60%
Epoch: 11 Learning rate: 0.0044262842543732585
Epoch [12/50], Training Loss: 0.5094, Training Accuracy: 82.2900 %
Test Epoch [12/50], Validation Loss: 0.4118, Validation Accuracy: 86.04%

Epoch: 12 Learning rate: 0.0043224229237103905
Epoch [13/50], Training Loss: 0.4836, Training Accuracy: 83.1900 %
Test Epoch [13/50], Validation Loss: 0.4059, Validation Accuracy: 85.92%
Epoch: 13 Learning rate: 0.004211369342086191
Epoch [14/50], Training Loss: 0.4585, Training Accuracy: 84.1280 %
Test Epoch [14/50], Validation Loss: 0.4719, Validation Accuracy: 84.16%
Epoch: 14 Learning rate: 0.004093561787251775
Epoch [15/50], Training Loss: 0.4419, Training Accuracy: 84.6680 %
Test Epoch [15/50], Validation Loss: 0.4321, Validation Accuracy: 85.26%
Epoch: 15 Learning rate: 0.003969465191804921
Epoch [16/50], Training Loss: 0.4162, Training Accuracy: 85.4600 %
Test Epoch [16/50], Validation Loss: 0.3843, Validation Accuracy: 86.77%
Epoch: 16 Learning rate: 0.0038395693083135155
Epoch [17/50], Training Loss: 0.4047, Training Accuracy: 85.9460 %
Test Epoch [17/50], Validation Loss: 0.3678, Validation Accuracy: 87.51%
Epoch: 17 Learning rate: 0.003704386776485917
Epoch [18/50], Training Loss: 0.3841, Training Accuracy: 86.6380 %
Test Epoch [18/50], Validation Loss: 0.3907, Validation Accuracy: 86.88%
Epoch: 18 Learning rate: 0.003564451100016223
Epoch [19/50], Training Loss: 0.3688, Training Accuracy: 87.0560 %
Test Epoch [19/50], Validation Loss: 0.4047, Validation Accuracy: 86.59%
Epoch: 19 Learning rate: 0.003420314541088931
Epoch [20/50], Training Loss: 0.3540, Training Accuracy: 87.6880 %
Test Epoch [20/50], Validation Loss: 0.3375, Validation Accuracy: 88.49%
Epoch: 20 Learning rate: 0.0032725459408523955
Epoch [21/50], Training Loss: 0.3410, Training Accuracy: 88.0680 %
Test Epoch [21/50], Validation Loss: 0.3565, Validation Accuracy: 88.33%
Epoch: 21 Learning rate: 0.0031217284744627
Epoch [22/50], Training Loss: 0.3224, Training Accuracy: 88.6340 %
Test Epoch [22/50], Validation Loss: 0.3796, Validation Accuracy: 87.34%
Epoch: 22 Learning rate: 0.002968457349557738
Epoch [23/50], Training Loss: 0.3102, Training Accuracy: 89.1260 %
Test Epoch [23/50], Validation Loss: 0.3621, Validation Accuracy: 88.32%
Epoch: 23 Learning rate: 0.0028133374572445924
Epoch [24/50], Training Loss: 0.2955, Training Accuracy: 89.6360 %
Test Epoch [24/50], Validation Loss: 0.3482, Validation Accuracy: 88.61%
Epoch: 24 Learning rate: 0.002656980984870685
Epoch [25/50], Training Loss: 0.2822, Training Accuracy: 90.1680 %

Test Epoch [25/50], Validation Loss: 0.3376, Validation Accuracy: 88.73%
Epoch: 25 Learning rate: 0.002500004999999999
Epoch [26/50], Training Loss: 0.2655, Training Accuracy: 90.6500 %
Test Epoch [26/50], Validation Loss: 0.3352, Validation Accuracy: 89.13%
Epoch: 26 Learning rate: 0.0023430290151293135
Epoch [27/50], Training Loss: 0.2599, Training Accuracy: 90.9260 %
Test Epoch [27/50], Validation Loss: 0.3208, Validation Accuracy: 89.70%
Epoch: 27 Learning rate: 0.0021866725427554068
Epoch [28/50], Training Loss: 0.2448, Training Accuracy: 91.4860 %
Test Epoch [28/50], Validation Loss: 0.3353, Validation Accuracy: 89.39%
Epoch: 28 Learning rate: 0.002031552650442261
Epoch [29/50], Training Loss: 0.2313, Training Accuracy: 91.9760 %
Test Epoch [29/50], Validation Loss: 0.3068, Validation Accuracy: 90.04%
Epoch: 29 Learning rate: 0.0018782815255372984
Epoch [30/50], Training Loss: 0.2155, Training Accuracy: 92.3560 %
Test Epoch [30/50], Validation Loss: 0.3307, Validation Accuracy: 89.71%
Epoch: 30 Learning rate: 0.0017274640591476039
Epoch [31/50], Training Loss: 0.2088, Training Accuracy: 92.7820 %
Test Epoch [31/50], Validation Loss: 0.3075, Validation Accuracy: 90.57%
Epoch: 31 Learning rate: 0.001579695458911069
Epoch [32/50], Training Loss: 0.1994, Training Accuracy: 93.0040 %
Test Epoch [32/50], Validation Loss: 0.3204, Validation Accuracy: 90.02%
Epoch: 32 Learning rate: 0.0014355588999837758
Epoch [33/50], Training Loss: 0.1861, Training Accuracy: 93.4720 %
Test Epoch [33/50], Validation Loss: 0.3466, Validation Accuracy: 89.84%
Epoch: 33 Learning rate: 0.0012956232235140817
Epoch [34/50], Training Loss: 0.1776, Training Accuracy: 93.8200 %
Test Epoch [34/50], Validation Loss: 0.3200, Validation Accuracy: 90.32%
Epoch: 34 Learning rate: 0.0011604406916864824
Epoch [35/50], Training Loss: 0.1644, Training Accuracy: 94.3000 %
Test Epoch [35/50], Validation Loss: 0.3059, Validation Accuracy: 90.88%
Epoch: 35 Learning rate: 0.0010305448081950786
Epoch [36/50], Training Loss: 0.1533, Training Accuracy: 94.7240 %
Test Epoch [36/50], Validation Loss: 0.3285, Validation Accuracy: 90.63%
Epoch: 36 Learning rate: 0.0009064482127482241
Epoch [37/50], Training Loss: 0.1479, Training Accuracy: 94.8620 %
Test Epoch [37/50], Validation Loss: 0.3326, Validation Accuracy: 90.37%

Epoch: 37 Learning rate: 0.0007886406579138076
Epoch [38/50], Training Loss: 0.1328, Training Accuracy: 95.4200 %
Test Epoch [38/50], Validation Loss: 0.3128, Validation Accuracy: 91.10%
Epoch: 38 Learning rate: 0.0006775870762896086
Epoch [39/50], Training Loss: 0.1267, Training Accuracy: 95.6260 %
Test Epoch [39/50], Validation Loss: 0.3037, Validation Accuracy: 91.40%
Epoch: 39 Learning rate: 0.0005737257456267409
Epoch [40/50], Training Loss: 0.1205, Training Accuracy: 95.8280 %
Test Epoch [40/50], Validation Loss: 0.3252, Validation Accuracy: 90.69%
Epoch: 40 Learning rate: 0.00047746655914760334
Epoch [41/50], Training Loss: 0.1119, Training Accuracy: 96.1400 %
Test Epoch [41/50], Validation Loss: 0.3200, Validation Accuracy: 90.88%
Epoch: 41 Learning rate: 0.00038918940788459027
Epoch [42/50], Training Loss: 0.1063, Training Accuracy: 96.2800 %
Test Epoch [42/50], Validation Loss: 0.3056, Validation Accuracy: 91.51%
Epoch: 42 Learning rate: 0.0003092426814237412
Epoch [43/50], Training Loss: 0.1000, Training Accuracy: 96.5680 %
Test Epoch [43/50], Validation Loss: 0.3167, Validation Accuracy: 91.27%
Epoch: 43 Learning rate: 0.00023794189297021392
Epoch [44/50], Training Loss: 0.0937, Training Accuracy: 96.7180 %
Test Epoch [44/50], Validation Loss: 0.3208, Validation Accuracy: 91.09%
Epoch: 44 Learning rate: 0.00017556843416180104
Epoch [45/50], Training Loss: 0.0948, Training Accuracy: 96.7520 %
Test Epoch [45/50], Validation Loss: 0.3189, Validation Accuracy: 91.31%
Epoch: 45 Learning rate: 0.0001223684645446976
Epoch [46/50], Training Loss: 0.0887, Training Accuracy: 96.9640 %
Test Epoch [46/50], Validation Loss: 0.3152, Validation Accuracy: 91.48%
Epoch: 46 Learning rate: 7.85519400942282e-05
Epoch [47/50], Training Loss: 0.0874, Training Accuracy: 97.0360 %
Test Epoch [47/50], Validation Loss: 0.3108, Validation Accuracy: 91.46%
Epoch: 47 Learning rate: 4.429178461453183e-05
Epoch [48/50], Training Loss: 0.0862, Training Accuracy: 97.0780 %
Test Epoch [48/50], Validation Loss: 0.3116, Validation Accuracy: 91.51%
Epoch: 48 Learning rate: 1.9723207287312153e-05
Epoch [49/50], Training Loss: 0.0867, Training Accuracy: 97.0460 %
Test Epoch [49/50], Validation Loss: 0.3121, Validation Accuracy: 91.58%
Epoch: 49 Learning rate: 4.943169062963242e-06
Epoch [50/50], Training Loss: 0.0830, Training Accuracy: 97.1860 %

Test Epoch [50/50], Validation Loss: 0.3121, Validation Accuracy: 91.53%

Epoch: 50 Learning rate: 1e-08


```
training_acc = []
for tensor in train_acc:
    training_acc.append(tensor.double().tolist())

# Create a list of epoch numbers
epochs = list(range(len(train_loss)))

# Plot the training loss and validation loss
plt.plot(epochs, train_loss, label='Training Loss')
plt.plot(epochs, v_loss, label='Validation Loss')

# Add a legend and axis labels
plt.legend()
plt.xlabel('Epoch')
plt.ylabel('Loss')

# Show the plot
plt.show()
```



```
# Plot the training accuracy and validation accuracy
plt.plot(epochs, training_acc, label='Training Accuracy')
```

```
plt.plot(epochs, v_acc, label='Validation Accuracy')
```

```
# Add a legend and axis labels
```


```
plt.legend()
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Accuracy')
```

```
# Show the plot
```

```
plt.show()
```



```
# Plot the training accuracy and validation accuracy
```

```
plt.plot(epochs, learn_rate, label='Learning Rate')
```

```
# Add a legend and axis labels
```


```
plt.legend()
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Learning rate')
```

```
# Show the plot
```

```
plt.show()
```



```
# Create a dictionary containing all the necessary information
checkpoint = {
    'model_state_dict': net.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'epoch': num_epochs,
    'accuracy': epoch_acc,
    'loss': epoch_loss
}

# Save the checkpoint to a file
torch.save(checkpoint, path)

# Load the saved data
checkpoint = torch.load(path)

# Extract the model state dictionary, optimizer state dictionary,
epoch number, accuracy, and loss
model_state_dict = checkpoint['model_state_dict']
optimizer_state_dict = checkpoint['optimizer_state_dict']
epoch = checkpoint['epoch']
accuracy = checkpoint['accuracy']
loss = checkpoint['loss']

# Load the saved state dictionaries into the model and optimizer
net.load_state_dict(model_state_dict)
```

```

optimizer.load_state_dict(optimizer_state_dict)

# Print all parameters of the model
for name, param in net.named_parameters():
    print(name, param.shape)
print(f"Loaded checkpoint from epoch {epoch} with accuracy {accuracy}")

conv1.weight torch.Size([32, 3, 3, 3])
conv1.bias torch.Size([32])
bn1.weight torch.Size([32])
bn1.bias torch.Size([32])
conv2.weight torch.Size([64, 32, 3, 3])
conv2.bias torch.Size([64])
bn2.weight torch.Size([64])
bn2.bias torch.Size([64])
conv3.weight torch.Size([128, 64, 3, 3])
conv3.bias torch.Size([128])
bn3.weight torch.Size([128])
bn3.bias torch.Size([128])
conv4.weight torch.Size([256, 128, 3, 3])
conv4.bias torch.Size([256])
bn4.weight torch.Size([256])
bn4.bias torch.Size([256])
conv5.weight torch.Size([512, 256, 3, 3])
conv5.bias torch.Size([512])
bn5.weight torch.Size([512])
bn5.bias torch.Size([512])
conv6.weight torch.Size([1024, 512, 3, 3])
conv6.bias torch.Size([1024])
bn6.weight torch.Size([1024])
bn6.bias torch.Size([1024])
fc1.weight torch.Size([512, 16384])
fc1.bias torch.Size([512])
fc2.weight torch.Size([10, 512])
fc2.bias torch.Size([10])
Loaded checkpoint from epoch 50 with accuracy 97.186

y_pred = []
y_true = []

# set model to eval mode and move to GPU if available
net.eval()
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
net.to(device)

# iterate over test data
for inputs, labels in test_loader:
    inputs, labels = inputs.to(device), labels.to(device)

```

```

with torch.no_grad():
    output = net(inputs) # Feed Network

    output = (torch.max(torch.exp(output), 1)
[1]).data.cpu().numpy()
    y_pred.extend(output) # Save Prediction

    labels = labels.data.cpu().numpy()
    y_true.extend(labels) # Save Truth

# constant for classes
classes = ('Airplane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog',
'horse', 'ship', 'truck')

# Build confusion matrix
cf_matrix = confusion_matrix(y_true, y_pred)
df_cm = pd.DataFrame(cf_matrix / np.sum(cf_matrix, axis=1)[:], None],
index = [i for i in classes],
                columns = [i for i in classes])
plt.figure(figsize = (12,10))
sn.heatmap(df_cm, annot=True)
plt.savefig('outputjugal.png')

```

